

Jupyter Random_Trees_CR Last Checkpoint: an hour ago (unsaved changes) Logout

File Edit View Insert Cell Kernel Help Trusted | Python 3

Run Code

Classification Using Random Forests (CR)

```
In [1]: import os
from matplotlib import pyplot as plt
from sklearn.feature_extraction import DictVectorizer
import pandas as pd
import numpy as np
from sklearn import datasets
from sklearn.tree import DecisionTreeRegressor
from sklearn import tree
```

Acquire data in an accessible format

```
In [2]: df = pd.read_csv(os.path.join(".", "Cleaned_Data", "chromatic.csv"))
df.head()
```

```
Out[2]:
```

	0	1	2	3	4	5	6	7	8	9	...	111	112	113	114	
0	7.161286	7.835325	2.911583	0.984049	-1.499548	-2.094097	0.578000	-1.205671	1.849122	-0.425568	...	-0.364194	-0.364194	-0.364194	-0.364194	-0
1	0.225763	-0.094169	-0.803646	0.497745	0.874036	0.290280	-0.077659	-0.887385	0.432082	-0.093963	...	0.938616	0.938616	0.938616	0.938616	0
2	-0.692525	-0.517801	-0.788035	1.214351	-0.907214	0.880213	0.406899	-0.694895	-0.901869	-1.701574	...	0.603755	0.603755	0.603755	0.603755	0
3	-0.735562	-0.684055	2.058215	0.716328	-0.011393	0.805396	1.497982	0.114752	0.692847	0.052377	...	0.187169	0.187169	0.187169	0.187169	0
4	0.570272	0.273157	-0.279214	0.083466	1.049331	-0.869295	-0.265858	-0.401676	-0.872639	1.147483	...	1.620715	1.620715	1.620715	1.620715	1

5 rows × 121 columns

Prepare data for the machine learning model Using the 'Sub_Region' Column as Input

```
In [3]: df.drop(columns=['Latitude', 'Longitude', 'Country', 'Region'])
df = df.drop(columns=['Latitude', 'Longitude', 'Country', 'Region'])
```

```
In [4]: Sub_Region_List = df['Sub_Region'].drop_duplicates()
print (Sub_Region_List)
```

```
0      South America
1      Western Africa
3      Eastern Africa
4      Northern Africa
8      Northern Europe
9      Central Asia
11     Southern Europe
12     Southern Asia
15     Eastern Europe
29     Western Asia
34     South-eastern Asia
42     Eastern Asia
53     Caribbean
76     Central America
97  Australia and New Zealand
Name: Sub_Region, dtype: object
```

Jupyter Random_Trees_CR Last Checkpoint: an hour ago (autosaved) Logout

File Edit View Insert Cell Kernel Help Trusted Python 3

One-Hot Encoding

```
In [6]: # Step 0: Reformat data
data = df.values
X = data[:, 0:115]
y = data[:, 116]
```

```
In [7]: from sklearn.preprocessing import LabelEncoder

# Step 1: Label-encode data set
label_encoder = LabelEncoder()
label_encoder.fit(y)
encoded_y = label_encoder.transform(y)
```

```
In [8]: from tensorflow.keras.utils import to_categorical

# Step 2: One-hot encoding
one_hot_y = to_categorical(encoded_y)
one_hot_y
```

```
Out[8]: array([[0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 1., 0.],
               [0., 0., 0., ..., 0., 1., 0.],
               ...,
               [0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
In [9]: for label, original_class in zip(encoded_y, y):
        print('Original Class: ' + str(original_class))
        print('Encoded Label: ' + str(label))
        print('-' * 15)
```

```
Original Class: Western Africa
Encoded Label: 13
-----
Original Class: Western Africa
Encoded Label: 13
-----
Original Class: Western Africa
Encoded Label: 13
-----
Original Class: Northern Europe
Encoded Label: 8
-----
Original Class: Central Asia
Encoded Label: 3
-----
Original Class: Central Asia
Encoded Label: 3
-----
Original Class: Southern Europe
Encoded Label: 12
-----
```

Separate The Data into Features & Targets

```
In [10]: target = one_hot_y
target_names = ["negative", "positive"]
```

```
In [11]: data = df.drop("Sub_Region", axis=1)
feature_names = data.columns
data.head()
```

```
Out[11]:
```

Jupyter Random_Trees_CR Last Checkpoint: an hour ago (autosaved) Python 3 Logout

File Edit View Insert Cell Kernel Help Trusted

Run Code

```

3 -0.735562 -0.684055 2.058215 0.718328 -0.011393 0.805396 1.497982 0.114752 0.892847 0.052377 ... 0.187169 0.187169 0.187169 0.187169 0
4 0.570272 0.273157 -0.279214 0.083458 1.049331 -0.869295 -0.265858 -0.401676 -0.872639 1.147483 ... 1.620715 1.620715 1.620715 1.620715 1
5 rows x 116 columns

```

Convert DataTypes For The Training & Testing Data Sets

In [12]: `df.dtypes`

Out[12]:

```

0      float64
1      float64
2      float64
3      float64
4      float64
...
112     float64
113     float64
114     float64
115     float64
Sub_Region    object
Length: 117, dtype: object

```

In [13]: `# Convert 'Sub_Region' Column to Float`

```

df["Sub_Region"] = pd.to_numeric(df.Sub_Region, errors='coerce')
df.dtypes

```

Out[13]:

```

0      float64
1      float64
2      float64
3      float64
4      float64
...
112     float64
113     float64
114     float64
115     float64
Sub_Region    float64
Length: 117, dtype: object

```

Train and split Model with random forest regression model from skicit-learn

In [14]:

```

# Import model
from sklearn.ensemble import RandomForestRegressor
# Instantiate model with 1000 decision trees
rf = RandomForestRegressor(n_estimators = 1000, random_state = 1234)
# Split Data Into Testing and Training Data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(data, target, random_state=1234)

```

Train & Split Model with Random Forest Classification Model

In [15]:

```

from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=200)
rf = rf.fit(X_train, y_train)
rf.score(X_test, y_test)

```

Out[15]: 0.88679245283018867

jupyter Random_Trees_CR Last Checkpoint: an hour ago (autosaved) Logout

File Edit View Insert Cell Kernel Help Trusted | Python 3

Run Code

Shape of all the Data

```
In [16]: print('X_train Shape:', X_train.shape)
print('y_train Shape:', y_train.shape)
print('X_test Shape:', X_test.shape)
print('y_test Shape:', y_test.shape)
```

```
X_train Shape: (794, 116)
y_train Shape: (794, 15)
X_test Shape: (265, 116)
y_test Shape: (265, 15)
```

Metrics and Scoring For Classification Model

```
In [17]: clf = tree.DecisionTreeClassifier()
clf = clf.fit(X_train, y_train)
clf.score(X_test, y_test)
```

```
Out[17]: 0.33584905660377357
```

```
In [18]: # Train the model on training data
rf = rf.fit(X_train, y_train)
rf.score(X_test, y_test)
```

```
Out[18]: 0.07924528301886792
```

Make Predictions & Calculate Errors

```
In [19]: # Use numpy to convert to arrays
import numpy as np
# Use the forest's predict method on the test data
predictions = rf.predict(X_test)
predictions
```

```
Out[19]: array([[0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
In [20]: # Use the forest's predict method on the test data
predictions = rf.predict(X_test)
# Calculate the absolute errors
errors = abs(predictions - y_test)
errors
```

```
Out[20]: array([[0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 1., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
In [21]: # Print out the mean absolute error (mae)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
```

```
Mean Absolute Error: 0.06 degrees.
```

```
In [22]: # Calculate mean absolute percentage error (MAPE)
# mape = 100 * (errors / y_test, )
# Calculate and display accuracy
# accuracy = 100 - np.mean(mape, )
```



Visualizing The Decision Tree in Regression Task

```
In [24]: # Fit the regressor, set max_depth = 3
regr = DecisionTreeRegressor(max_depth=3, random_state=1234)
model = regr.fit(X, one_hot_y)
```

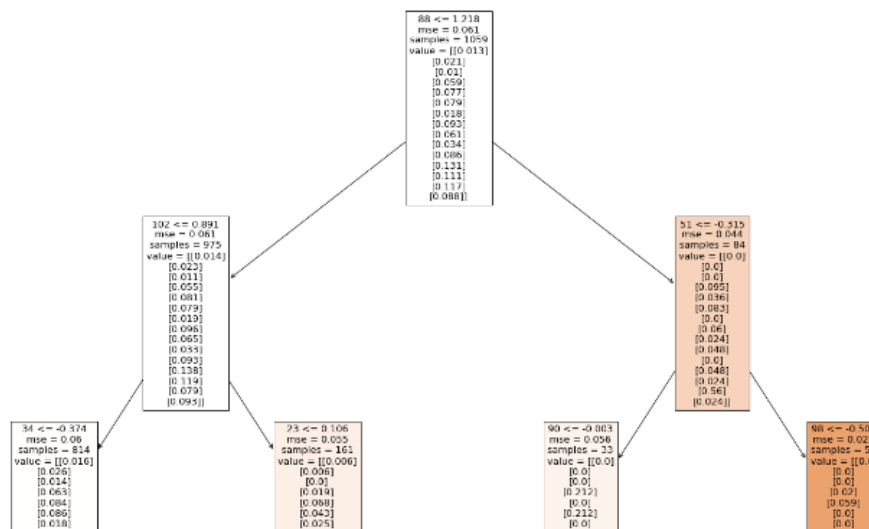
```
In [25]: text_representation = tree.export_text(regr)
print(text_representation)
```

```

--- feature_88 <= 1.22
    --- feature_102 <= 0.89
        --- feature_34 <= -0.37
            --- value: [0.01, 0.01, 0.00, 0.05, 0.03, 0.13, 0.01, 0.10, 0.04, 0.02, 0.04, 0.24, 0.12, 0.03, 0.18]
            --- feature_34 > -0.37
                --- value: [0.02, 0.04, 0.02, 0.07, 0.12, 0.05, 0.03, 0.12, 0.04, 0.05, 0.05, 0.05, 0.12, 0.12, 0.06]
        --- feature_102 > 0.89
            --- feature_23 <= 0.11
                --- value: [0.00, 0.00, 0.00, 0.02, 0.07, 0.05, 0.02, 0.03, 0.02, 0.03, 0.46, 0.09, 0.13, 0.08, 0.02]
                --- feature_23 > 0.11
                    --- value: [0.02, 0.02, 0.00, 0.02, 0.06, 0.04, 0.04, 0.02, 0.56, 0.00, 0.06, 0.04, 0.08, 0.04, 0.00]
            --- feature_88 > 1.22
                --- feature_51 <= -0.31
                    --- feature_90 <= -0.00
                        --- value: [0.00, 0.00, 0.00, 0.30, 0.00, 0.04, 0.00, 0.09, 0.04, 0.04, 0.00, 0.04, 0.04, 0.30, 0.09]
                        --- feature_90 > -0.00
                            --- value: [0.00, 0.00, 0.00, 0.00, 0.00, 0.60, 0.00, 0.00, 0.10, 0.20, 0.00, 0.10, 0.00, 0.00, 0.00]
                    --- feature_51 > -0.31
                        --- feature_98 <= -0.50
                            --- value: [0.00, 0.00, 0.00, 0.00, 0.10, 0.00, 0.00, 0.30, 0.00, 0.10, 0.00, 0.10, 0.10, 0.30, 0.00]
                            --- feature_98 > -0.50
                                --- value: [0.00, 0.00, 0.00, 0.02, 0.05, 0.00, 0.00, 0.00, 0.00, 0.00, 0.00, 0.02, 0.00, 0.90, 0.00]

```

```
In [26]: # Note that color of the leaf corresponds to the predicted value.
fig = plt.figure(figsize=(25,20))
_ = tree.plot_tree(regr, feature_names=data.columns, filled=True)
fig.savefig("static/images/random_trees_DF")
```



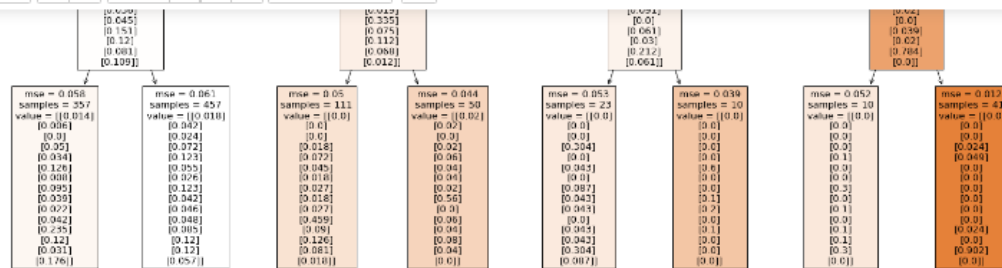
File Edit View Insert Cell Kernel Help

Trusted

Python 3 ○

 Run

 Code



Random Trees Data Structure and Visualization

```
In [27]: # Total Count of Nodes
```

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
n_nodes = rf.estimators_[0].tree_.node_count
n_nodes
```

```
Out[27]: 503
```

```
In [28]: # Visualize (5) Random Trees
```

```
# Note ALSO that color of the Leaf corresponds to the predicted value.
fn=data.columns
cn=one_hot_y
```

```
fig, axes = plt.subplots(nrows = 1, ncols = 5, figsize = (12,2), dpi=720)
for index in range(0, 5):
    tree.plot_tree(rf.estimators_[index],
                   feature_names = fn,
                   class_names=cn,
                   filled = True,
                   ax = axes[index]);
```

```
axes[index].set_title('Estimator: ' + str(index), fontsize = 11)
```

```
fig.savefig("static/images/rt_streets-CR.png")
```



In []: